

## ASSIGNMENT-4

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style("darkgrid")
```

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn import metrics
```

```
df=pd.read_csv('/content/abalone.csv')
```

```
df.head(10)
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395
5	I	0.425	0.300	0.095	0.3515	0.1410	0.0775
6	F	0.530	0.415	0.150	0.7775	0.2370	0.1415
7	F	0.545	0.425	0.125	0.7680	0.2940	0.1495
8	M	0.475	0.370	0.125	0.5095	0.2165	0.1125
9	F	0.550	0.440	0.150	0.8945	0.3145	0.1510

	Shell weight	Rings
0	0.150	15
1	0.070	7
2	0.210	9
3	0.155	10
4	0.055	7
5	0.120	8
6	0.330	20
7	0.260	16
8	0.165	9
9	0.320	19

```
df.shape
```

```
(4177, 9)
```

```
df.describe()
```

	Length	Diameter	Height	Whole weight	Shucked weight \
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367
std	0.120093	0.099240	0.041827	0.490389	0.221963
min	0.075000	0.055000	0.000000	0.002000	0.001000
25%	0.450000	0.350000	0.115000	0.441500	0.186000
50%	0.545000	0.425000	0.140000	0.799500	0.336000
75%	0.615000	0.480000	0.165000	1.153000	0.502000
max	0.815000	0.650000	1.130000	2.825500	1.488000

	Viscera weight	Shell weight	Rings
count	4177.000000	4177.000000	4177.000000
mean	0.180594	0.238831	9.933684
std	0.109614	0.139203	3.224169
min	0.000500	0.001500	1.000000
25%	0.093500	0.130000	8.000000
50%	0.171000	0.234000	9.000000
75%	0.253000	0.329000	11.000000
max	0.760000	1.005000	29.000000

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 4177 entries, 0 to 4176
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Sex	4177 non-null	object
1	Length	4177 non-null	float64
2	Diameter	4177 non-null	float64
3	Height	4177 non-null	float64
4	Whole weight	4177 non-null	float64
5	Shucked weight	4177 non-null	float64
6	Viscera weight	4177 non-null	float64
7	Shell weight	4177 non-null	float64
8	Rings	4177 non-null	int64

```
dtypes: float64(7), int64(1), object(1)
```

```
memory usage: 293.8+ KB
```

```
df['age']=df['Rings']+1.5
```

```
df=df.drop('Rings', axis = 1)
```

```
df.head(10)
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight \
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395

5	I	0.425	0.300	0.095	0.3515	0.1410	0.0775
6	F	0.530	0.415	0.150	0.7775	0.2370	0.1415
7	F	0.545	0.425	0.125	0.7680	0.2940	0.1495
8	M	0.475	0.370	0.125	0.5095	0.2165	0.1125
9	F	0.550	0.440	0.150	0.8945	0.3145	0.1510

	Shell weight	age
0	0.150	16.5
1	0.070	8.5
2	0.210	10.5
3	0.155	11.5
4	0.055	8.5
5	0.120	9.5
6	0.330	21.5
7	0.260	17.5
8	0.165	10.5
9	0.320	20.5

```
df.isnull().sum()
```

```
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Shucked weight 0
Viscera weight 0
Shell weight 0
age          0
dtype: int64
```

```
df.columns
```


```
Index(['Sex', 'Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
      'Viscera weight', 'Shell weight', 'age'],
      dtype='object')
```

```
sns.distplot(df['Length'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-
level function with similar flexibility) or `histplot` (an axes-level
function for histograms).
```

```
warnings.warn(msg, FutureWarning)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6285125e10>
```




```
sns.boxplot(df.Sex,df.Height)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:
```

```
FutureWarning: Pass the following variables as keyword args: x, y. From  
version 0.12, the only valid positional argument will be `data`, and passing  
other arguments without an explicit keyword will result in an error or  
misinterpretation.
```


```
FutureWarning
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f62847d9810>
```




```
sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x7f62842be8d0>
```




```
df = pd.get_dummies(df)
dummy_df = df

var = 'Viscera weight'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```



```
df.drop(df[(df['Viscera weight'] > 0.5) &
           (df['age'] < 20)].index, inplace = True)
df.drop(df[(df['Viscera weight'] < 0.5) & (
df['age'] > 25)].index, inplace = True)


var = 'Shell weight'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```



```

df.drop(df[(df['Shell weight'] > 0.6) & (df['age'] < 25)].index, inplace =
True)
df.drop(df[(df['Shell weight'] < 0.8) & (df['age'] > 25)].index, inplace =
True)
var = 'Shucked weight'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)

```




```


df.drop(df[(df['Whole weight'] >= 2.5) & (df['age'] < 25)].index, inplace =
True)
df.drop(df[(df['Whole weight'] < 2.5) & (df['age'] > 25)].index, inplace =
True)
var = 'Diameter'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)

```






```
df.drop(df[(df['Diameter'] < 0.1) & (df['age'] < 5)].index, inplace = True)
df.drop(df[(df['Diameter'] < 0.6) & (df['age'] > 25)].index, inplace = True)
df.drop(df[(df['Diameter'] >= 0.6) & (df['age'] < 25)].index, inplace = True)
var = 'Height'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```



```
df.drop(df[(df['Height'] > 0.4) & (df['age'] < 15)].index, inplace = True)
df.drop(df[(df['Height'] < 0.4) & (df['age'] > 25)].index, inplace = True)
var = 'Length'
```

```
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```



```
numerical_features = df.select_dtypes(include = [np.number]).columns
categorical_features = df.select_dtypes(include = [np.object]).columns
```

/usr/local/lib/python3.7/dist-packages/ipykernel\_launcher.py:2:  
 DeprecationWarning: `np.object` is a deprecated alias for the builtin  
 `object`. To silence this warning, use `object` by itself. Doing this will  
 not modify any behavior and is safe.  
 Deprecated in NumPy 1.20; for more details and guidance:  
<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
numerical_features
```


```
Index(['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',  
      'Viscera weight', 'Shell weight', 'age', 'Sex_F', 'Sex_I', 'Sex_M'],  
      dtype='object')
```

```
categorical_features
```

```
Index([], dtype='object')
```

```
plt.figure(figsize = (24,7))
sns.heatmap(df[numerical_features].corr(),annot = True)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f627e059c90>
```



```
df.columns
```

```
Index(['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
      'Viscera weight', 'Shell weight', 'age', 'Sex_F', 'Sex_I', 'Sex_M'],
      dtype='object')
```

```
X = df.drop('age', axis = 1)
y = df['age']
```

## LINEAR REGRESSION

```
from sklearn.feature_selection import SelectKBest
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score
standardScale = StandardScaler()
standardScale.fit_transform(X)
```

```
selectkBest = SelectKBest()
X_new = selectkBest.fit_transform(X, y)
```

```
X_train, X_test, y_train, y_test = train_test_split(X_new, y, test_size =
0.25)
```

```
lr = LinearRegression()
lr.fit(X_train, y_train)
```

```
LinearRegression()
```

```
y_train_pred = lr.predict(X_train)
y_test_pred = lr.predict(X_test)
```

```
from sklearn.metrics import mean_absolute_error, mean_squared_error
s = mean_squared_error(y_train, y_train_pred)
print('Mean Squared Error of training set :%2f'%s)
```

```
p = mean_squared_error(y_test, y_test_pred)
print('Mean Squared Error of testing set :%2f'%p)
```

Mean Squared Error of training set :4.458678  
Mean Squared Error of testing set :4.748683

```
from sklearn.metrics import r2_score  
s = r2_score(y_train, y_train_pred)  
print('R2 Score of training set:%.2f'%s)
```

```
p = r2_score(y_test, y_test_pred)  
print('R2 Score of testing set:%.2f'%p)
```

R2 Score of training set:0.53  
R2 Score of testing set:0.53